

Program 6:

AIM: Develop a Pig Latin Scripts to sort, group, join, Project and filter the data

Source Code:

1. Sample Data

Let's assume we have two datasets:

- employees.txt (Employee ID, Name, Age, Department ID, Salary)
- departments.txt (Department ID, Department Name)

employees.txt (stored in HDFS at /user/cloudera/employees.txt)

101,John,30,1,50000

102,Sam,28,2,60000

103,Anna,32,1,75000

104,David,29,3,62000

105,Lily,27,2,58000

departments.txt (stored in HDFS at /user/cloudera/departments.txt)

1,HR

2,Finance

3,IT

Pig Latin Script

Save the script as employee_analysis.pig and execute it in Cloudera.

-- Load the employees dataset

```
employees = LOAD 'hdfs://localhost:9000/user/cloudera/employees.txt'
```

```
USING PigStorage(',')
```

```
AS (emp_id:int, name:chararray, age:int, dept_id:int, salary:int);
```

-- Load the departments dataset

```
departments = LOAD 'hdfs://localhost:9000/user/cloudera/departments.txt'
```

```
USING PigStorage(',')
```

```
AS (dept_id:int, dept_name:chararray);
```

-- 1. FILTER: Select employees with age greater than 28

```
filtered_employees = FILTER employees BY age > 28;
```

-- 2. PROJECT: Select only emp_id, name, and salary columns

```
projected_employees = FOREACH filtered_employees GENERATE emp_id, name, salary;
```

-- 3. SORT: Order employees by salary in descending order

```
sorted_employees = ORDER projected_employees BY salary DESC;
```

```
-- 4. GROUP: Group employees by department ID
grouped_by_department = GROUP employees BY dept_id;

-- 5. JOIN: Join employees with department names using dept_id
joined_data = JOIN employees BY dept_id, departments BY dept_id;

-- STORE results in HDFS
STORE sorted_employees INTO 'hdfs://localhost:9000/user/cloudera/output/sorted_employees' USING
PigStorage(',');

STORE grouped_by_department INTO
'hdfs://localhost:9000/user/cloudera/output/grouped_by_department' USING PigStorage(',');
STORE joined_data INTO 'hdfs://localhost:9000/user/cloudera/output/joined_data' USING PigStorage(',');

-- DISPLAY the results on the screen
DUMP sorted_employees;
DUMP grouped_by_department;
DUMP joined_data;
```

Upload the data in HDFS

```
hdfs dfs -mkdir -p /user/cloudera
hdfs dfs -put employees.txt /user/cloudera/
hdfs dfs -put departments.txt /user/cloudera/
```

Run the Script

```
pig -x mapreduce employee_analysis.pig
```

output commands

```
hdfs dfs -cat /user/cloudera/output/sorted_employees/part-r-00000
hdfs dfs -cat /user/cloudera/output/grouped_by_department/part-r-00000
hdfs dfs -cat /user/cloudera/output/joined_data/part-r-00000
```

OUTPUT :Sorted Employee by Salary

```
103,Anna,75000
104,David,62000
102,Sam,60000
101,John,50000
```

Grouped Employees by Departments

```
(1, {(101, John, 30, 1, 50000), (103, Anna, 32, 1, 75000)})  
(2, {(102, Sam, 28, 2, 60000), (105, Lily, 27, 2, 58000)})  
(3, {(104, David, 29, 3, 62000)})
```

Joined Employees with Departments

```
(101, John, 30, 1, 50000, 1, HR)  
(102, Sam, 28, 2, 60000, 2, Finance)  
(103, Anna, 32, 1, 75000, 1, HR)  
(104, David, 29, 3, 62000, 3, IT)  
(105, Lily, 27, 2, 58000, 2, Finance)
```

Program 7:

AIM: Use HIVE to create, alter, and drop databases, tables, views, functions and indexes

Create a Database:

```
CREATE DATABASE employee_db;
```

Output:

```
OK  
Time taken: 0.234 seconds
```

Use the database:

```
USE employee_db;
```

Output:

```
OK  
Time taken: 0.123 seconds
```

Alter a Database:

```
ALTER DATABASE employee_db SET DBPROPERTIES ('Owner'='Admin');
```

```
OK  
Time taken: 0.134 seconds
```

Drop a Data Base

```
DROP DATABASE employee_db CASCADE;
```

```
OK  
Time taken: 0.321 seconds
```

Create a Table

```
CREATE TABLE employees (  
    emp_id INT,  
    name STRING,  
    age INT,  
    dept_id INT,  
    salary FLOAT  
)
```

```
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE;
```

Output:

```
OK
Time taken: 0.567 seconds
```

```
LOAD DATA INPATH '/user/cloudera/employees.txt' INTO TABLE employees;
```

Output:

```
Loading data to table employee_db.employees
OK
Time taken: 1.543 seconds
```

Alter a Table

Add a New Column

```
ALTER TABLE employees ADD COLUMNS (email STRING);
```

Output:

```
OK
Time taken: 0.234 seconds
```

Rename:

```
ALTER TABLE employees RENAME TO employees_new;
```

Drop Table:

```
DROP TABLE employees_new;
```

Create a View

```
CREATE VIEW high_salary_employees AS
SELECT emp_id, name, salary
FROM employees
WHERE salary > 50000;
```

Alter a View:

```
ALTER VIEW high_salary_employees AS  
SELECT emp_id, name, age, salary  
FROM employees  
WHERE salary > 60000;
```

Drop View:

```
DROP VIEW high_salary_employees;
```

Create a Function

Add a JAR file containing a Java-based UDF:

```
ADD JAR /user/cloudera/custom_udf.jar;
```

```
CREATE FUNCTION to_upper AS 'com.example.hiveudf.ToUpperUDF';
```

```
Added /user/cloudera/custom_udf.jar to class path  
OK  
Time taken: 0.568 seconds
```

Use the Function

```
SELECT to_upper(name) FROM employees;
```

Output:

```
JOHN  
SAM  
ANNA  
DAVID  
LILY  
OK  
Time taken: 0.345 seconds
```

Drop Function

```
DROP FUNCTION to_upper;
```

Create an Index

```
CREATE INDEX emp_dept_idx  
ON TABLE employees (dept_id)  
AS 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler'  
WITH DEFERRED REBUILD;
```

Output:

```
OK  
Time taken: 0.765 seconds
```

Rebuild the Index:

```
ALTER INDEX emp_dept_idx ON employees REBUILD;
```

Drop Index

```
DROP INDEX emp_dept_idx ON employees;
```

Check all tables in the current database:

```
SHOW TABLES;
```

```
default  
employee_db  
OK  
Time taken: 0.167 seconds
```

Check all tables in the current database:

```
SHOW TABLES;
```

```
employees  
employees_new  
OK  
Time taken: 0.145 seconds
```

DESCRIBE employees;

```
emp_id      int
name        string
age         int
dept_id     int
salary      float
email       string
```

OK

Time taken: 0.234 seconds

Display the table data

SELECT * FROM employees LIMIT 5;

```
101  John  30  1  50000.0  NULL
102  Sam   28  2  60000.0  NULL
103  Anna  32  1  75000.0  NULL
104  David 29  3  62000.0  NULL
105  Lily  27  2  58000.0  NULL
```

OK

Time taken: 0.459 seconds

Program 8:

AIM: Implement word count program in Hadoop and spark

Source Code;

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;
import java.util.StringTokenizer;

public class WordCount {

    public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException
        {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

```

public static class IntSumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,
    InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Compile and Package the Java Code

```

javac -classpath `hadoop classpath` -d . WordCount.java
jar cf wc.jar WordCount*.class

```

Run the Hadoop Word Count Job

Upload the input file to HDFS:

```

hdfs dfs -mkdir -p /user/cloudera/input
hdfs dfs -put sample.txt /user/cloudera/input/

```

Run the Map Reduce job

```
hadoop jar wc.jar WordCount /user/cloudera/input /user/cloudera/output
```

View the Output:

```
hdfs dfs -cat /user/cloudera/output/part-r-00000
```

```
Hadoop      2
Hello       3
MapReduce   1
Spark       2
World       2
```

Word Count in Apache Spark (PySpark)

1.Prepare the Input File

Upload the Sample.txt

```
hdfs dfs -mkdir -p /user/cloudera/input
```

```
hdfs dfs -put sample.txt /user/cloudera/input/
```

sample input File:

Hello Hadoop

Hello Spark

Hello World

Spark is fast

Hadoop is slow

2. Create a new script file using nano or vi:

```
nano wordcount.py
```

3.code

```
from pyspark.sql import SparkSession
```

```
# Initialize Spark Session
```

```
spark = SparkSession.builder.appName("WordCount").getOrCreate()
```

```
# Read text file from HDFS
text_file = spark.sparkContext.textFile("hdfs://localhost:9000/user/cloudera/input/sample.txt")

# Process data
word_counts = (text_file
               .flatMap(lambda line: line.split(" ")) # Split lines into words
               .map(lambda word: (word, 1))           # Map each word to (word, 1)
               .reduceByKey(lambda a, b: a + b))      # Reduce by key (word) and sum counts

# Save the output to HDFS
word_counts.saveAsTextFile("hdfs://localhost:9000/user/cloudera/output_spark")

# Print results
for word, count in word_counts.collect():
    print(f"{word}: {count}")

# Stop Spark Session
spark.stop()
```

3. Execute the Script

```
spark-submit wordcount.py
```

4. View the output:

```
hdfs dfs -ls /user/cloudera/output_spark
```

```
hdfs dfs -cat /user/cloudera/output_spark/part-00000
```

```
Hello 3
Hadoop 2
Spark 2
World 1
is 2
fast 1
slow 1
```

Program 9:

AIM: Use CDH (Cloudera Distribution for Hadoop) and HUE (Hadoop User Interface) to analyze the data and generate reports for sample data sets

Steps and Source Code:

1.Start Cloudera services:

```
sudo service cloudera-scm-server start
```

```
sudo service cloudera-scm-agent start
```

Check the status:

```
sudo service --status-all | grep cloudera
```

2. Access HUE Web Interface

Open your browser and navigate to:

<http://localhost:8888>

Login using HUE credentials:

Username: cloudera

Password: cloudera

3. Upload Sample Dataset to HDFS

Example Dataset: Employee Data (employees.csv)

id,name,department,salary

1,John,IT,70000

2,Alice,HR,60000

3,Bob,IT,75000

4,Charlie,Finance,80000

5,David,HR,62000

6,Eva,IT,72000

7,Frank,Finance,81000

8,Grace,HR,65000

Upload employees.csv to HDFS using HUE

1. **Navigate to HUE → File Browser → HDFS**
2. **Create a new directory** /user/cloudera/data
3. Click **Upload** → Select employees.csv → Upload it

Or through terminal

```
hdfs dfs -mkdir -p /user/cloudera/data
```

```
hdfs dfs -put employees.csv /user/cloudera/data/
```

Create a Hive Table in HUE

Open HUE → Click Query Editors → Select Hive

4. Create a Hive table for the dataset:

```
CREATE DATABASE IF NOT EXISTS company;
```

```
USE company;
```

```
CREATE TABLE employees (
```

```
  id INT,
```

```
  name STRING,
```

```
  department STRING,
```

```
  salary INT
```

```
)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY ','
```

```
STORED AS TEXTFILE;
```

Load data into the Hive table:

```
LOAD DATA INPATH '/user/cloudera/data/employees.csv' INTO TABLE employees;
```

Verify the data:

```
SELECT * FROM employees;
```

id	name	department	salary
1	John	IT	70000
2	Alice	HR	60000
3	Bob	IT	75000
4	Charlie	Finance	80000
5	David	HR	62000
6	Eva	IT	72000
7	Frank	Finance	81000
8	Grace	HR	65000

Data Analysis Using Hive Queries in HUE

Find the Highest Salary in Each Department

```
SELECT department, MAX(salary) AS highest_salary  
FROM employees  
GROUP BY department;
```

Output:

Department	Highest Salary
IT	75000
HR	65000
Finance	81000

Get Employees with Salary Greater Than 65000;

Output:

Name	Department	Salary
John	IT	70000
Bob	IT	75000
Eva	IT	72000
Charlie	Finance	80000
Frank	Finance	81000

Generate Reports in HUE

Step 1: Export Query Results

1. Run any of the above SQL queries in HUE Query Editor
2. Click Export → Choose format (CSV, Excel, JSON)
3. Download the report

Step 2: Create HUE Dashboard for Visualization

1. Open HUE → Click Dashboard
2. Click Create New Dashboard
3. Click Add Widget → Select Chart Type (Bar Chart, Pie Chart, etc.)

Enter Query → Example for Employee Count:

```
SELECT department, COUNT(*) AS employee_count FROM employees GROUP BY department;
```

Click **"Run Query"** → The visualization will be generated

Department	Employee Count
IT	3
HR	3
Finance	2